CLIP learns by matching images with their captions; we keep only the image tower

discarded in our detector

similarity of every image with every caption

$u_1 \cdot v_1$	$u_1 \cdot v_2$	$u_1 \cdot v_3$
$u_2 \cdot v_1$	$u_2 \cdot v_2$	$u_2 \cdot v_3$
$u_3 \cdot v_1$	$u_3 \cdot v_2$	$u_3 \cdot v_3$

training pulls matching pairs (red diagonal) together and pushes others apart

Layout is schematic.