The greedy one-to-one rule on a toy example (made-up numbers)

step 1: real r_1 picks

	f_1	f_2	f_3	f_4
r_1	0.2	0.9	0.4	0.1
r_2	0.3	0.8	0.7	0.2
r_3	0.6	0.5	0.9	0.3

r_1: best is f_2 (0.9).

step 2: real r_2 picks

	f_1	f_2 ×	f_3	f_4
r_1	0.2	0.9	0.4	0.1
r_2	0.3	0.8	0.7	0.2
r_3	0.6	0.5	0.9	0.3

r_2: best is f_2, but it is used (×).
Take the next best, f_3 (0.7).

step 3: real r_3 picks

	f_1	f_2 ×	f_3 ×	f_4
r_1	0.2	0.9	0.4	0.1
r_2	0.3	0.8	0.7	0.2
r_3	0.6	0.5	0.9	0.3

r_3: f_3 and f_2 are used (×).
Take the best left, f_1 (0.6).

Rows are real images, columns are fakes, cells are cosine similarities. Each fake is used at most once.