Stochastic weight averaging: keep a running mean of the weights

weights at the end of each epoch (the first update copies; the average mixes epochs 13–19)

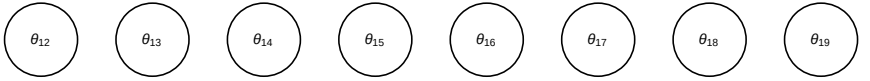

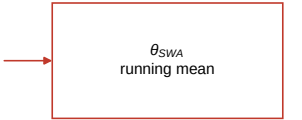

$$\theta_{SWA}^{(t)} = \theta_{SWA}^{(t-1)} + \frac{\theta_t - \theta_{SWA}^{(t-1)}}{n_t + 1}$$

three deployable models are kept: best regular (a single epoch's weights), best SWA (average at its best validation epoch), final SWA (average of all)

the learning rate is lowered to 10^{-5} with a ten-epoch cosine schedule while averaging